

Documenting PH levels

(Non-technical report)

Jun Yan, Ritesh Lohiya, Justin Herman

May 9, 2019

Problem

New regulations are requiring us to understand our manufacturing process, the predictive factors and be able to report to them our predictive model of PH.

Process

At ABC beverage we take our products and government compliance issues extremely seriously. As such, my data science team, has conducted a thorough analysis of the PH levels in our mineral water. You can see the detailed approach my team took [here](#). The purpose of this report is to document some of the findings in a digestible manner for both the executives and government compliance officials.

We collected over 2800 samples of our beverages and documented over 33 different measurements of our manufacturing processes. Below I present to you a chart as 10 examples of our processes are recorded and measured

Brand Code	Carb Volume	Fill Ounces	PC Volume	Carb Pressure	Carb Temp	PSC	PSC Fill	PSC CO2	Mnf Flow	Carb Pressure1	Fill Pressure	Hyd Pressure1	Hyd Pressure2	Hyd Pressure3	Hyd Pressure4
B	5.34	23.97	0.26	68.2	141.2	0.10	0.26	0.04	-100	118.8	46.0	0	NA	NA	118
A	5.43	24.01	0.24	68.4	139.6	0.12	0.22	0.04	-100	121.6	46.0	0	NA	NA	106
B	5.29	24.06	0.26	70.8	144.8	0.09	0.34	0.16	-100	120.2	46.0	0	NA	NA	82
A	5.44	24.01	0.29	63.0	132.6	NA	0.42	0.04	-100	115.2	46.4	0	0	0	92
A	5.49	24.31	0.11	67.2	136.8	0.03	0.16	0.12	-100	118.4	45.8	0	0	0	92

Filler Level	Filler Speed	Filler Temperature	Usage cont	Carb Flow	Carb Density	MFR	Balling	Pressure Vacuum	Oxygen Filler	Bowl Setpoint	Pressure Setpoint	Air Pressurer	Alch Rel	Carb Rel	Balling Lvl	PH
121.2	4002	66.0	16.18	2932	0.88	725.0	1.40	-4.0	0.02	120	46.4	142.6	6.58	5.32	1.48	8.36
118.6	3986	67.6	19.90	3144	0.92	726.8	1.50	-4.0	0.03	120	46.8	143.0	6.56	5.30	1.56	8.26
120.0	4020	67.0	17.76	2914	1.58	735.0	3.14	-3.8	0.02	120	46.6	142.0	7.66	5.84	3.28	8.94
117.8	4012	65.6	17.42	3062	1.54	730.6	3.04	-4.4	0.03	120	46.0	146.2	7.14	5.42	3.04	8.24
118.6	4010	65.6	17.68	3054	1.54	722.8	3.04	-4.4	0.03	120	46.0	146.2	7.14	5.44	3.04	8.26

Figure 1-First Five Observations

Our main target (variable we want to explore) is PH. Looking at PH above, it seems like it is mostly in the 8-9 range. As you can see on figure 2 below, this is an accurate assumption. Our average PH is around 8.5, and the target is normally distributed. Normally distributed variables are great for predictive purposes, as prediction is a mathematically intensive process and normally distributed variables meet many of the assumptions needed to run complex models




Figure 2-Boxplot and Histogram of PH

As you are aware, due to government regulations we need to give a report on our entire manufacturing process and how it effects our PH. Our goal is to use the data in the tables above, to predict future levels of PH.

Exploratory Data Analysis (EDA)

Through exploratory data analysis, we can better understand our data. Healthy predictions are dependent upon healthy data. The first step we take in our EDA is understanding how our predictors (our measurements) correlate with our target.

	PH		PH
Carb Volume	0.07	Filler Speed	0.08
Fill Ounces	-0.12	Temperature	-0.18
PC Volume	0.1	Usage cont	-0.36
Carb Pressure	0.08	Carb Flow	0.23
Carb Temp	0.03	Density	0.1
PSC	-0.07	MFR	0.05
PSC Fill	0.05	Balling	0.08
PSC CO2	-0.09	Pressure Vacuum	0.22
Mnf Flow	-0.46	PH	1
Carb Pressure1	-0.12	Oxygen Filler	0.16
Fill Pressure	-0.32	Bowl Setpoint	0.36
Hyd Pressure1	-0.05	Pressure Setpoint	-0.31
Hyd Pressure2	-0.22	Air Pressurer	
Hyd Pressure3	-0.27	Alch Rel	0.17
Hyd Pressure4	-0.17	Carb Rel	0.2
Filler Level	0.35	Balling Lvl	0.11

Figure 3- Measuring PH Correlations

Understanding correlation is more of an art than a science, but correlations that matter less, are faded out in the above graph. We can assume that MNF flow, Fill. Pressure, Filler Level, , Usage.Cont, Bowl. Setpoint and Pressure. Setpoint are all very impactful on our PH levels. We should keep this in mind when looking at the results of our models.

Internal Notes on Data Collection (EDA)

My team was not in charge of the data collection. In terms of missing values, the data collection team did an excellent job. There were only 4 observations where our PH wasn't recorded in the entire dataset. Below you can see a printout of our dataset, where the red values are our missing values. The chart of the left indicates that out MFR measurement(predictors) had about 8% of the data missing. Typically, this would be on the threshold for discarding a predictor. Overall, the other predictors are well within statistical limits. We decided not to drop any of this data, and we ran KNN imputation to fill in the missing data. This method can predict how to best fill in our missing data based upon KNN classification. Please see the Technical report for further details on how we tested imputation methods and their application in our models.




Figure 4- Missing Data

Potential Issues with Data Collection

(This section of the report is for internal executives only)

We attempted to contact the data collection team, but we were unable to receive a satisfactory answer to some of the questions we have found within the data. On the following page in figure 5, Bowl. Setpoint seems to be measured in intervals of 10 from 70-120, and then the variable begins to take on several different even integer values. Pressure set point seems to be measured by integers, and then there are 3 observations between 46 and 48. PSC.CO2 seems to be suffering from a clear rounding issue.

Observations of .1999999996 and of .02000000005 are very likely to both represent a measurement of .2. We want to document these issues, so that when future data collection occurs, we can prevent such issues. In terms of our current modeling approach, we kept the data as is, and believe it is unlikely that correcting the data would have resulted in materially different results

Bowl.Setpoint	Freq	PSC.CO2	Freq
70	99	0	108
80	96	0.01999999999999996	325
90	434	0.02000000000000005	288
100	112	0.03999999999999991	37
110	437	0.04	624
120	1307	0.05999999999999996	283
122	1	0.06000000000000005	219
126	10	0.07999999999999992	23
130	51	0.08	234
134	2	0.09999999999999996	79
140	20	0.10000000000000001	65
		0.11999999999999999	10
Pressure.Setpoint	Freq	0.12	72
44	96	0.14	30
46	1322	0.14000000000000001	19
46.4	1	0.15999999999999999	3
46.6	1	0.16	36
46.8	1	0.18	12
48	125	0.18000000000000001	8
50	1002	0.19999999999999999	2
52	11	0.2	16
		0.22	22
		0.24	17

Figure 5- Irregularities in Data

Modeling

We attempted several different modeling solutions. All our models are fivefold cross validated and tuned for best performance. Cross validation statistically tests our model against different samples of the data so that we can be confident the model will be predictive of unforeseen data. In terms of Linear models, we ran a PLS and an Elastic Net model. We use RMSE as our error measurement, but others were reported.

Linear models suffer from an inability to generalize nonlinear trends. These linear models allow us more flexibility to fit our models to the data. As seen above in figure 6, the PLS RMSE error is around.137 and

the elastic net is around .140. These models have been tuned to select their best hyperparameters. Hyperparameters allow us to determine optimal points in which our model has minimized bias (prediction accuracy on our dataset), while minimizing variance (ability to predict unforeseen data)

We tested two nonlinear models as well (KNN(RMSE=.127), SVM (RMSE= .135)).

We conclude our modeling process by testing out tree-based models. Random Forest and XgBoost were chosen. XgBoost does not actually need data imputation, so we tested models with and without imputation. Preliminary results indicated that imputation models performed better (please see technical writeup). Our Tree models produced the best results of all models. Below you can see out Random Forest, which reaches our lowest error rate of .117

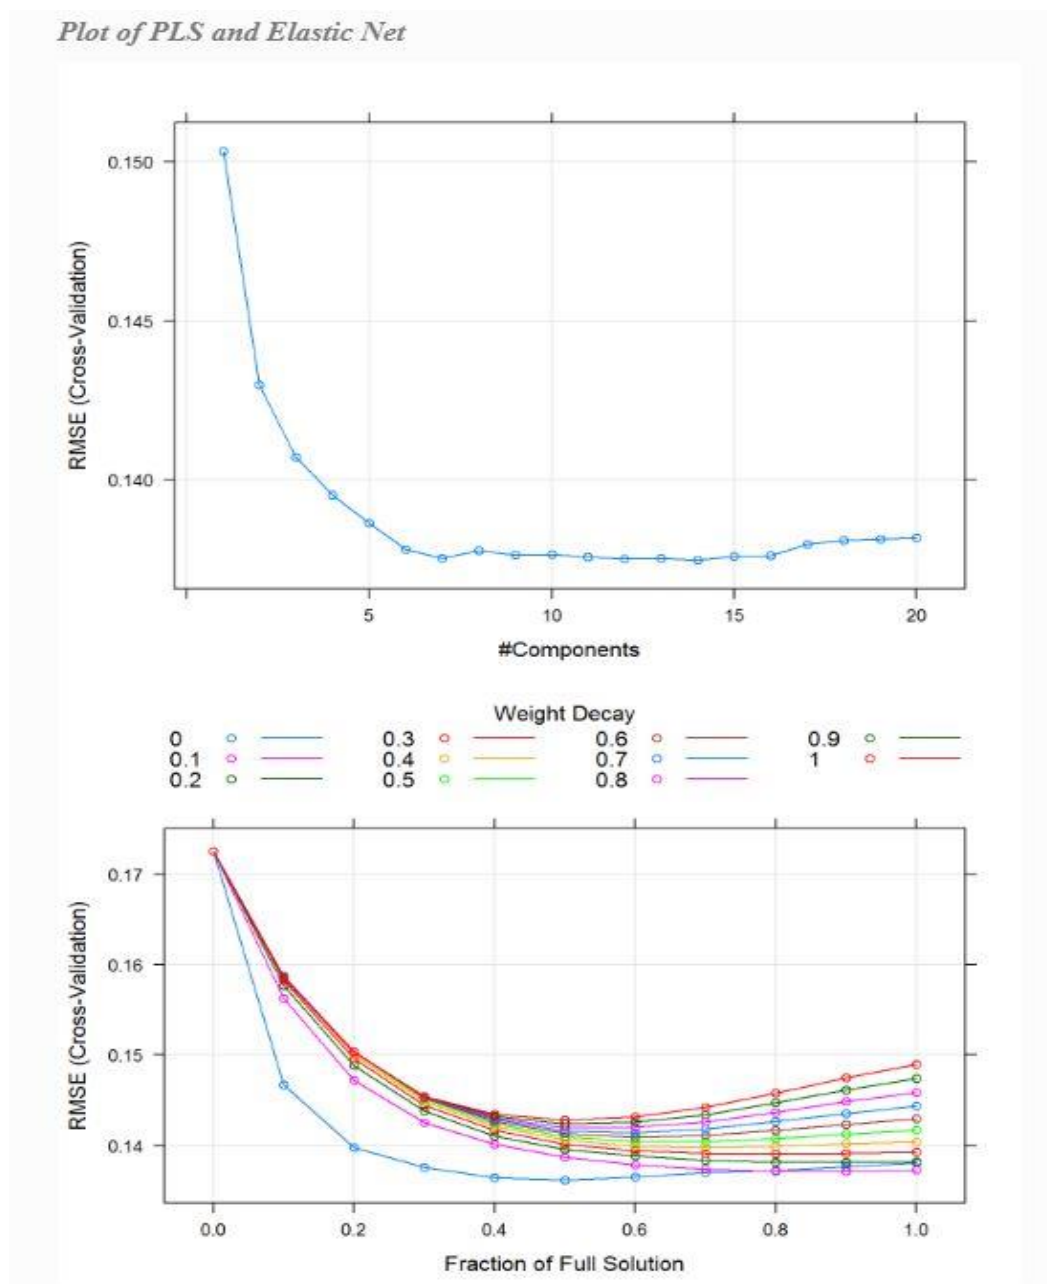



Figure 6- PLS (top) & Elastic Net(bottom) CV tuning

Plot KNN Model



Plot SVM Model

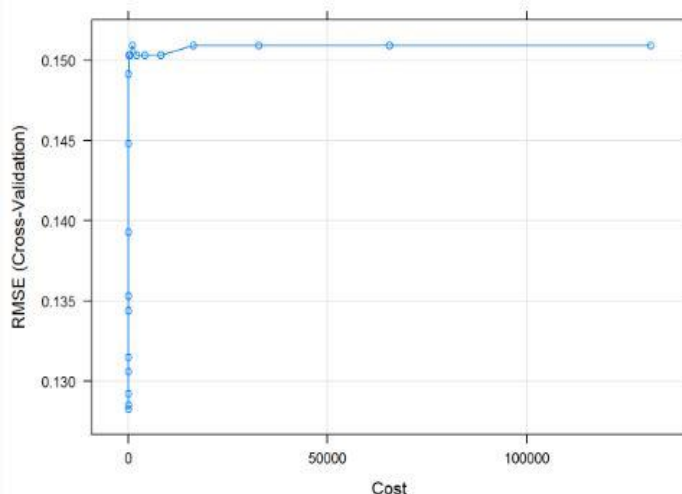


Figure 7-Nonlinear Model Tuning-KNN(left) SVM(right)

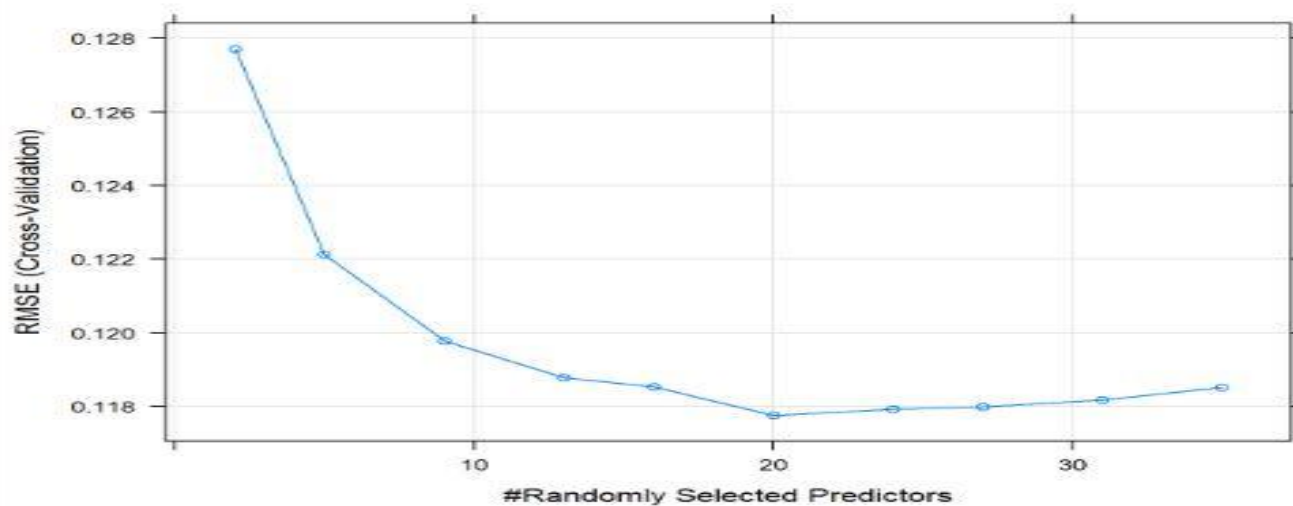


Figure 8-Random Forest

The XgBoost model (RMSE= .116) doesn't produce a graph, but as you can see from the table summarizing all the results from our models (figure 9), XgBoost produces the strongest cross validated performance.

Models: PLS, ENet, KNN, SVM, RF, XGB

Number of resamples: 4

MAE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
PLS	0.10401017	0.10573777	0.10715156	0.10662121	0.10803500	0.10817153
ENet	0.10325386	0.10502571	0.10631413	0.10584920	0.10713763	0.10751469
KNN	0.10357093	0.10386052	0.10401850	0.10396019	0.10411817	0.10423283
SVM	0.09480799	0.09557060	0.09584557	0.09591785	0.09619282	0.09717227
RF	0.08744219	0.08787527	0.08848266	0.08871182	0.08931921	0.09043978
XGB	0.08697277	0.08713229	0.08730968	0.08735522	0.08753261	0.08782876

NA's

PLS	0
ENet	0
KNN	0
SVM	0
RF	0
XGB	0

RMSE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
PLS	0.1356567	0.1371824	0.1377113	0.1374039	0.1379328	0.1385362	0
ENet	0.1329212	0.1356141	0.1362803	0.1357745	0.1364407	0.1366161	0
KNN	0.1344314	0.1345692	0.1349825	0.1357495	0.1361629	0.1386017	0
SVM	0.1259120	0.1265990	0.1273402	0.1278324	0.1285736	0.1307371	0
RF	0.1165398	0.1172607	0.1177933	0.1177503	0.1182830	0.1188748	0
XGB	0.1156593	0.1157390	0.1161317	0.1164846	0.1168773	0.1180158	0

Rsqared

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
PLS	0.3625243	0.3656883	0.3683503	0.3689665	0.3716286	0.3766411	0
ENet	0.3700989	0.3761005	0.3796023	0.3791235	0.3826253	0.3871907	0
KNN	0.3790417	0.3836021	0.3902105	0.3906420	0.3972504	0.4031051	0
SVM	0.4293764	0.4446522	0.4546945	0.4522908	0.4623331	0.4703980	0
RF	0.5379883	0.5399803	0.5430877	0.5458798	0.5489872	0.5593555	0
XGB	0.5351413	0.5400408	0.5417958	0.5459440	0.5476990	0.5650431	0

Figure 9-Overall Model performance

Variable Importance

The results of our model are rather encouraging. Our RMSE is extremely low relative to our target variable. We are unsure as of now what regulatory burden will be placed on our company, however, we can confidentially predict our PH levels given the data. An RMSE of .116 means that on average the error squared of any prediction is within .116 of PH. Assuming we need to eventually adjust the PH levels of our mineral water, we have confidentially identified the predictors which most influence our PH. Below you can see the variable importance from all the models we ran.

Rank	pls	rf	xgbTree	enet	knn	svmRadial
1	MnfFlow	MnfFlow	MnfFlow	MnfFlow	MnfFlow	MnfFlow
2	BrandCodeC	BrandCodeC	Usagecont	MFR	MFR	MFR
3	BowlSetpoint	PressureVacuum	BrandCodeC	BowlSetpoint	BowlSetpoint	BowlSetpoint
4	Usagecont	OxygenFiller	OxygenFiller	FillerLevel	FillerLevel	FillerLevel
5	FillerLevel	BallingLvl	AlchRel	PressureSetpoint	PressureSetpoint	PressureSetpoint
6	HydPressure3	AirPressurer	Temperature	Usagecont	Usagecont	Usagecont
7	PressureSetpoint	AlchRel	PressureVacuum	CarbFlow	CarbFlow	CarbFlow
8	FillPressure	CarbRel	FillerSpeed	BrandCodeC	BrandCodeC	BrandCodeC
9	BrandCodeB	Usagecont	CarbPressure1	HydPressure3	HydPressure3	HydPressure3
10	Temperature	Temperature	CarbRel	FillPressure	FillPressure	FillPressure
11	HydPressure2	CarbFlow	AirPressurer	PressureVacuum	PressureVacuum	PressureVacuum
12	PressureVacuum	FillerSpeed	BowlSetpoint	HydPressure2	HydPressure2	HydPressure2
13	CarbPressure1	HydPressure3	CarbFlow	CarbRel	CarbRel	CarbRel
14	OxygenFiller	CarbPressure1	Balling	HydPressure4	HydPressure4	HydPressure4
15	CarbFlow	Density	BallingLvl	Temperature	Temperature	Temperature
16	BrandCodeD	BowlSetpoint	FillerLevel	BrandCodeD	BrandCodeD	BrandCodeD
17	AlchRel	HydPressure2	Density	OxygenFiller	OxygenFiller	OxygenFiller
18	CarbRel	Balling	PCVolume	FillerSpeed	FillerSpeed	FillerSpeed
19	BrandCodeA	PCVolume	MFR	AlchRel	AlchRel	AlchRel
20	BallingLvl	FillPressure	HydPressure2	CarbPressure1	CarbPressure1	CarbPressure1

Figure 10- Variable Importance

Variable importance tells us the predictors which most influenced our PH. You can see that universally MNF flow was selected as our most important predictor. Bowl. Setpoint, Transcode, Filler Level, and Pressure Setpoint, all seem to be selected often in terms of importance in our models. Earlier on in the report, from simply looking at the correlations of PH to our predictors we speculated that (MNF flow, Fill. Pressure, Filler Level, Usage.Cont, Bowl. Setpoint and Pressure. Setpoint) would be important. This is useful in many ways.

Highly correlated variables being used in our models often is confirmation that our model is working correctly, as we assumed these variables should be important to our model. In the future we can develop simpler linear models using the predictors we know are likely to have predictive value. We can also better understand how to manipulate our future manufacturing process assuming new regulations require so. MNF flow is our minimum night flow. We know that with a correlation of $-.46$, as we increase MNF, our PH tends to decrease. Manipulation of this one measurement may in fact be enough to produce results to meet future regulatory burden

Conclusion

XGBoost has produced an extremely accurate model with a cross-validated RMSE of $.116$. We suggest that this report be used to initialize further experiments. We would like to test the effects of our importance predictors in a live setting. We also suggest that data collection errors that we noted be addressed in our future data collection. We are confident that by calibrating predictors in our manufacturing process to certain levels, Beverage ABC will be able to manipulate our PH levels to be prepared for whatever future regulatory burden is placed on us in the future.

